Quantum Circuit Optimilization Using Monte Carlo Tree Search and Reinforcement Learning

Author: Liza Darwesh
BSc. Thesis

Supervisors:
Ariana Torres, SURF
Damian Podareanu, SURF
Dick Heinhuis, HvA

Hogeschool van Amsterdam

Abstract

...partial- couple of sentences about motivation / task

Contents

List of Figures

Chapter 1

Introduction

1.1 Motivation

It is fascinating to think about the fact that the computing power of a military computer from 50 years ago, that was the size of an entire room, are in today's regular computers. However, even with the phenomenal growth we made in technology, there remain challenges which regular computers cannot seem to solve efficiently or effectively.

Quantum computers are expected to make a breakthrough in chemical and biological engineering through the discovery and manipulation of molecules; encryption for cyber security; the processing of very large quantities to aid in artificial intelligence; and the pricing of complex assets in finance [4]. Today's quantum computers perform tasks that exceed the capabilities of today's regular computers, but the world is not there yet.

1.2 Problem

Think of a spinning coin with information: either heads or tails. Eventually the coin will stop spinning and land on one of the sides. This is the same for a quantum system. A quantum computer consists of quantum bits (qubit) that have the same spinning effect as a spinning coin. In the case of a qubit, it is not heads or tails, but 0 or 1. One of the barriers to programming a quantum computer is that a qubit does not keep spinning for long. This can cause information to be lost, making some input algorithms have an unreliable output.

The algorithms that are written for a quantum computer contain a noise, which causes qubits to stop spinning faster. The way algorithms can be operated is by writing them into quantum gates. They are the same as the regular logic gates, but with quantum phenomena, which will be explained later in the paper. Mathematical algorithms are transformed into circuits of operations. These circuits are written in Quantum Assembly, a language that the quantum computer understands. Each circuit has a defined number of logical qubits. These qubits are what the gates are interacting with. In the moment of processing, the logical qubits in the circuit will then be mapped into the topology of the quantum computer, which is a connectivity architecture of the qubits inside the quantum

computer, also refereed as physical qubits. Another barrier of programming a quantum computer is matching the logical qubits with the topology, in a way that it is beneficial for the circuit. When programming a circuit, one must satisfy the topology. For instance, say two qubits need to interact with each other through a gate. If these qubits are not connected to each other, this operation will be inoperable, which means the whole circuit cannot be executed. So there needs to be a way to make circuits operable while satisfying the connectivity constraints of the quantum computer. A way to do this is by adding a so called SWAP-gate in the circuit, that causes the qubits to switch. This switch ensures that the qubits come closer to each other or even become connected. The problem with the SWAP-gate is that it is not a known logic gate for the quantum computer. This will be described later in the paper. Shortly said, there needs to be an optimization in placing these SWAP-gates to ensure that the information is reliable.

1.3   Organisation

The research of this thesis is done for the organisation SURF. A cooperative association of Dutch educational and research institutions. Universities, hospitals and research institutes work together within SURF to purchase or develop the best possible digital services. The company owns supercomputers which are available to complete high performance processes. Within the company there are multiple departments. The Innovation lab is one of these departments, which includes quantum innovation, machine learning, and high performance computing. Which will be useful for this research.

1.4   State-of-the-art

There are already a number of studies done in this area, for example a state-of-the-art research was conducted by Pozzi, et al. [5], using Reinforcement Learning and something called Quantum Annealing to optimize the number of SWAP-gates in a circuit. The research was based on a random logical qubit allocation on the physical qubits. When creating an initial qubit placement, the number of SWAP-gates can be reduced. Another recently published study was done by Sinha, et al [6], where the approach managed to minimize the number of SWAP-gates for a random qubit allocation using the Monte Carlo Tree Search in combination with Reinforcement Learning.

1.5   Research

Both papers discussed the need to look into initial qubit placement, which will also ensure that the SWAP gates are minimized in addition to using Reinforcement Learning. With this in mind, research will be conducted on quantum circuit optimization using the Monte Carlo Tree Search in combination with Reinforcement Learning. An initial qubit placement will be used for this.

This paper answers the main question: 'How can quantum circuit routing be optimized using Reinforcement Learning applied on Monte Carlo Tree Search?' This question will be answered by means of the following sub-questions:

- Why are there problems with quantum circuit routing?

- What has already been done to successfully accomplish quantum circuit routing?

- How will an initial qubit placement procedure help with the minimization of SWAP-gates?

- How does the Monte Carlo Tree Search work on optimizing the number of SWAP-gates?

This research paper is organised as follows. Section 2 offers a description of basic knowledge about quantum computing phenomenon and techniques needed for understanding the main question. Section 3 demonstrates what research has already been done by others. Section 4 describes the sub-problem initial qubit placement and how to solve it. Section 5 provides an overview of the experimental setup, which is how the Monte Carlo Tree Search is put together with the Reinforcement Learning. Section 6 shows the results of what the reinforcement learning model delivered. Finally, in Section 7 and 8 states the discussion about the research that is performed and the final conclusion.

Chapter 2

Quantum Circuit Routing

It is recognized that quantum computers can provide revolutionary developments. This chapter gives answer to the question 'what are the problems with quantum circuit routing' by first explaining what a quantum computer is and how it can be programmed. After that, an explanation will be given on the problem with an understanding of what needs to be optimized.

2.1 The Qubit

Figure 2.1: Regular bit and quantum bit illustration

A quantum computer can be seen as a regular computer that stores information and performs operations using quantum mechanics. A regular computer stores all its information such as numbers, text, and images, in series of 0's and 1's. These units are called bits. The way a quantum computer stores information is through the use of quantum bits (qubits). The difference between a bit and a qubit is that a qubit can carry not only the information 0 or 1, but also 0 and 1 on top of each other. This phenomenon is called superposition. A well-known example of superposition is the Schrödinger's cat example, where a hypothetical cat is illustrated in a closed box with a bottle of poison. It is not certain whether the cat is alive or dead unless it is checked. Before checking, the cat finds itself in a situation where it has a 50% chance of being dead or alive. This is the same with a qubit, there is no certainty whether it will hold 0 or 1, only a 50% probability that it will be one of the two. See Figure 2.1 for an illustration of a bit and a qubit.

2.2   Quantum Circuits

A regular computer operates by means of directions (algorithm) that are provided in the form of a script. This script then tells the computer processor what calculations have to be done to arrive at a certain output. These scripts can be written in different programming languages and the computer translates this into binary language, to control the logic gates inside the processor. The logic gates provide the calculations necessary for the operations. This principle, of programming gates that do calculations to perform a certain operation, work the same in a quantum computer. A script is written that determines which quantum gates should operate. An example of a gate that will be discussed frequently in this thesis is the Controlled NOT gate (CNOT-gate), see figure 2.2.

Figure 2.2: Example of a logic gate in a regular computer (left), example of a Quantum gate (right)

This script then creates a circuit consisting of the sequence of the quantum gates that were defined in the script. The algorithm that has been programmed is thus represented as a circuit with quantum gates. An example of a circuit can be seen in figure 2.3. The circuit illustrates between which qubits operations have to be performed. The circuit of figure 2.3 illustrates a four qubit circuit consisting of CNOT-gates with each interacting with two qubits. Usually, qubits are represented in Dirac notation [7], but in this thesis that is not relevant and the qubits will be represented with a capital Q.

Figure 2.3: A quantum circuit extracted from a programmed script

Each horizontal line in the circuit indicates the 'path' of a qubit, each path can have a quantum gate through which the qubit has to pass. Ultimately, the paths can be read to obtain a final result. As mentioned earlier in the example

of Schrödinger's cat, there is a 50% chance that the cat is alive or dead unless the box is opened. Reading or measuring a circuit forces the qubits to choose a state, either 0 or 1. The moment it chooses one of the two, the qubit collapses to a regular bit with the one value. This means that the output of the circuit will represent an ordinary series of bits, just like in a regular computer.

2.3   Topology

The step after feeding a quantum circuit to the quantum computer is the process in which the defined qubits in the circuit (logical qubits) are allocated to the qubits in the hardware (physical qubits). The logical qubits are, as it were, assigned a location on the hardware. The form in which these qubits can be allocated depends on the connectivity architecture of the physical qubits. This means that the qubits in the hardware are attached to each other in a certain form. This architecture is also known as the topology of the quantum computer.

There are different types of topologies. For example, figure 2.4b shows the topology of Rigetti's Aspen-4 quantum computer with a linear structure, or figure 2.4e shows the topology of Google's Sycamore, which consists of more connected qubits in a grid structure.

Figure 2.4: Quantum topologies of big tech companies' quantum computers. Source: [1]

These connectivities determine which qubits are allowed to interact with each other. Imagine the topology of figure 2.5 and the circuit of figure 2.3.

Figure 2.5: Example of a linear topology.

There are several ways to map the circuit's logical qubits into the computer's physical qubits, see figure 2.6. It can be done in sequence, randomly, but also with an initial placement, where the allocation is based on the qubits that need to interact in the circuit.

Figure 2.6: Different combinations of qubit mapping.

After mapping the logical qubits, the operations in the circuit will be performed. The moment interactions have to be performed in the circuit between qubits that have no connection in the topology, the operation cannot be performed, after which the entire circuit is rejected by the quantum computer. This appears to be a common problem when writing quantum algorithms. One way to make the circuit perform on a quantum computer is to add SWAP-gates to the circuit 2.7.

Figure 2.7: A SWAP-gate and its decomposition.

The SWAP gate will ensure that the qubits, that interact with the SWAP-gate, will switch positions in the quantum computer. This means that the qubits that have to operate with each other have a connection in the topology. This can be done over the entire circuit, until all operations are executable hardware-based.

A SWAP-gate is not a standard gate known to the quantum computer. A SWAP-gate can be simulated by placing three CNOT gates in a row, as shown in figure 2.7. Because there are actually three extra gates added per SWAP gate, this will ensure that the circuit will have a greater depth.

2.4 Decoherence

Qubits can be compared to a spinning coin. Eventually the coin will stop spinning and land on heads or tails. A qubit is also spinning between 0 and 1, but will also

eventually stop spinning, after which it will end up at 0 or 1. This phenomenon of breaking out of the superposition, after which it has to choose a state is called decoherence. The lifespan of a qubit differs per hardware composition of the quantum computer. Based on the temperature and network in the quantum chip, it can be ensured that the lifespan of a qubit is longer. The longest achieved so far by a qubit is 50 seconds, but the most common is somewhere between 15-120µ seconds [8]. Because qubits decohere, it causes parts of information to be lost. As a result, the output information will be unreliable. For this reason, it is important to set up a circuit that is small in depth, so that it can be quickly executed by the quantum computer while the qubits last. This is a tricky problem when a large algorithm has to be executed, where SWAP-gates are essential.

Because it is important to have a circuit that can be executed on a quantum computer, but also has to be small in depth, it is necessary to look for a way to optimize quantum circuits. A number of approaches have already been undertaken for this purpose. These can be found in the next chapter.

Chapter 3

Previously Researched Work

This chapter will provide a picture of what has been done before to optimize the quantum circuit routing and what has been recommended by these researchers for further research.

3.1 Cambridge's $t\,|ket\rangle$'s Routing Algorithm

Several approaches have been taken to optimize the circuit routing problem. They have all shown a state-of-the-art application. A well-known one is the use of the routing algorithm from Cambridge's Quantum Software Development Kit (Q-SDK) $t\,|ket\rangle$ [9]. The $t\,|ket\rangle$ Q-SDK is designed to maximize the performance of quantum algorithms when run on quantum computing hardware and accelerate the development of quantum computing applications across multiple industrial sectors. Cambridge students construed a solution to the circuit routing problem implemented with $t\,|ket\rangle$ [10]. They believe that this heuristic method in $t\,|ket\rangle$ matches or outperforms the results of other circuit mapping systems in terms of depth and total gates of the compiled circuit, and has a much shorter run time allowing larger circuits to be routed. They envision a quantum computer as a graph in which nodes are qubits and edges are the allowed interactions between the qubits. Their new architecture-agnostic methodology for mapping quantum circuits to real-world quantum computers with limited qubit connectivity present empirical results demonstrating the effectiveness of this method.


Figure 3.1: Example of a quantum circuit containing two-qubit gates acting on four qubits, Q0, Q1, Q2 and Q3. This circuit has six timesteps, each with gates acting on disjoint sets of qubits.

A procedure was defined by $t\,|ket\rangle$ for the initial mapping of logical qubits into

physical qubits. They repeat all the gate operations to construct a graph whose vertices are qubits. The gate operations in a circuit can be separated by means of boundaries, see Figure 3.1. Each separation is referred to as a timestep. At each time step, they add an edge to the graph between the interacting qubits present in the timestep. Decoupled qubits in the graph correspond either to qubits that have no interaction at all, or to qubits whose first interacts with a qubit whose first two interacts with others. These decoupled qubits are not included in the initial placement. They are added later in the routing procedure. This graph will eventually have nodes with multiple edges. The edge with the most edges is allocated first. The way this is done, will be explained in chapter 4. This ensures that most of the gates in the first two timesteps can be applied without any SWAP-gates. If the initial mapping cannot be completed respectfully to the topology, consideration should be given to the placement of SWAP-gates.

3.2 SWAP-based BidiREctional heuristic search algorithm

The next approach is inspired by the paper written by Gushu Li et al. [3] where they describe, a SWAP-based BidiREctional heuristic search algorithm, named SABRE, proposed to solve the qubit mapping problem. IBM's Q-SDK, called Qiskit, created a function part of a transpiler library based on the algorithms described in their paper. Transpilation is the process of rewriting a particular input circuit to match the topology of a particular quantum device, and/or to optimize the circuit for execution on today's noisy quantum systems. Most circuits must undergo a series of transformations that make them compatible with a particular device and optimize them to reduce the effects of noise on the resulting results. Rewriting quantum circuits to meet hardware limitations and optimizing for performance can be far from trivial. The flow of logic in the rewrite tool chain does not have to be linear and can often have iterative subloops, conditional branches, and other complex behaviour. That being said, the basic building blocks follow the structure given in Figure 3.2.
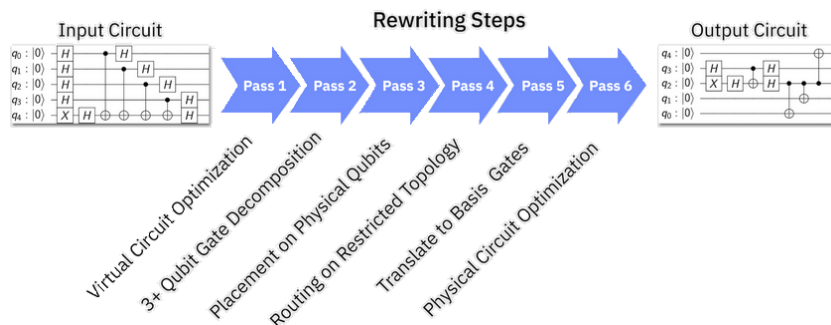

Figure 3.2: IBM's Transpilation Processes for Circuit Optimization, Source: [2]

The algorithm aims to minimize the number of SWAPs inserted and the circuit depth. This algorithm starts with an initial layout of logical qubits on physical qubits and repeats the circuit until all gates are exhausted, inserting SWAPs along

the way. This algorithm starts with Pass 3 from the process shown in Figure 3.2 and goes to Pass 1 after the initialization of the qubit placement. In each iteration, it is first checked whether there are gates in the front time step that can be applied directly. If so, it will apply them and check for the next timestep. Otherwise, it will try to search for SWAP-gates, insert the SWAP-gates, and update the mapping. The search for SWAP-gates is limited in that Qiskit considers only physical qubits close to those qubits in the previous time step.

Observing that many exhaustive search attempts may be redundant and that effective initial placement must begin with the qubits in the gates to be executed, Li et al. designed an optimized SWAP-based heuristic search scheme in SABRE with significantly reduced search space. They present a novel search technique in SABRE to naturally generate a initial mapping through traversing a reverse circuit, in which more consideration is given to those gates at the beginning of the circuit without completely ignoring the rest of the circuit, see Figure 3.3 .

Figure 3.3: Initial Mapping Update Using Reverse Traversal Technique, Source: [3]

In addition, they introduced a heuristic cost function for evaluating overlapping SWAP-gates, so SABRE tended to select non-overlapping SWAP ports. This optimization generates several hardware compatible circuits with a trade-off between circuit depth and number of gates.

3.3 Reinforcement Learning Approaches

The most popular approach for circuit routing optimization is through deep learning with reinforcement. In the article written by Matteo Pozzi et al. [5], they describe a qubit routing procedure using a modified version of the deep Q-learning model. Where they benchmark with the most advanced quantum compilers currently available (Qiskit and t|ket), on both random and realistic circuits. This research was a follow-up to the original research by Herbert and Sengupta [11]. Introducing the concept of reinforcement learning (RL) for the qubit routing problem. The way the RL model learns is through feeding the agent (neural network model) the current state and the reward that comes along with it. That way can the model can decide what action to take next, see Figure 3.4. The concept of RL will be described in more details in chapter 5.

The state that is represented to the agent, defined by the writers, contains the following information: the initial location of the qubit, its interaction with another logical qubit, and if it is already possible to carry out a CNOT-gate immediately. The agent has to select some SWAP-gates of the remaining qubits that are not involved in the CNOT-gates. By training a model, they try to carry

Figure 3.4: An illustration of the RL process

out the routing. The state representation goes through a function that computes a feature representation, which condenses the state in a fixed length vector of distances (how many qubits does the model have to go through). They feed the vector through the neural network, which outputs a single continues number which is the Q-value (quality of the current state) in the context of Q-Learning, a learning reinforcement learning technique, based on the quality of state and next state. They use the concept of simulated annealing to select a bunch of SWAPs to add the action in sequences on possible SWAP-gates and evaluates the quality. This process is depicted in Figure 3.5. Simulated Annealing is a simulation of the thermodynamics Annealing concept. Simply explained, annealing is where the system ends up in better solution, in a landscape of solutions, through a fast decrease in temperature of the system. The simulated annealing encourages the exploration of the state space and decreases the temperature, so the system ends up exploring less and less, until it ends up only adding a SWAP-gate when it is benefiting the quality.
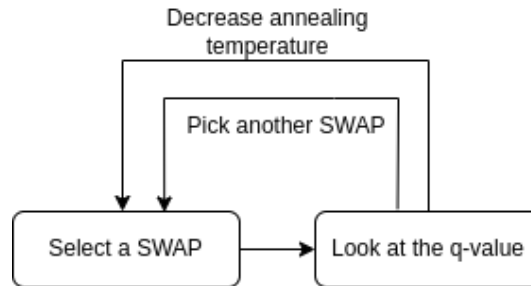
Figure 3.5: Q-Learning process with Simulated Annealing

They mentioned that in order to achieve this approach, it is necessary to alter the conventional RL framework to allow combinatorial action space, where they mention numerical results that are included to demonstrate the advantage of the RL-trained qubit routing policy over using a sorting network. Their approach turned out to be one of the best approaches, with results that turned out to be better than the approaches used by IBM and Cambridge.

3.4 Graph Neural Network with Monte Carlo Tree Search

The last approach is the use of the Monte Carlo Tree Search (MCTS) guided by a Graph neural network. In the paper written by Sinha et al. [6] they mention that the depth of the transformed quantum circuits is minimized by utilizing the MCTS to perform qubit routing by making it both construct each action and search over the space of all actions. It is aided in performing these tasks by a Graph neural network that evaluates the value and action probabilities for each state. The way MCTS is briefly shown in Figure 3.6. The way MCTS works is more thoroughly explained in chapter 5.
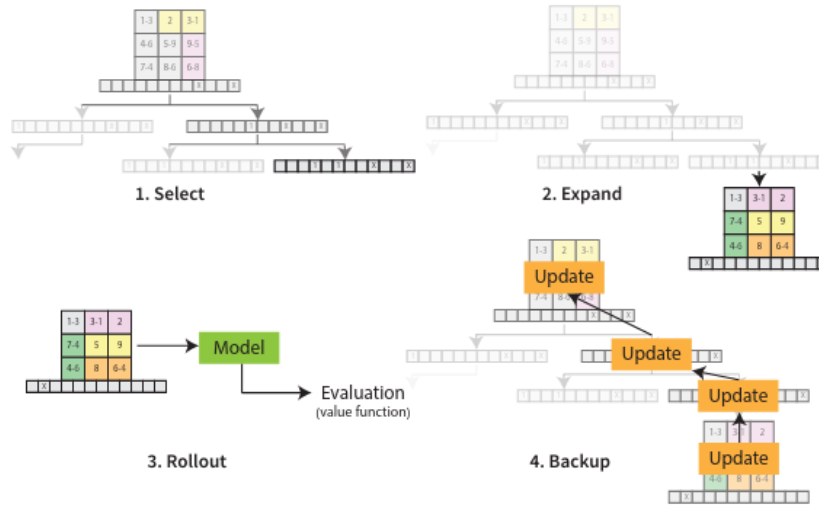
Figure 3.6: Iteration of a Monte Carlo tree search: (1) select - recursively choosing a node in the search tree for exploration using a selection criteria, (2) expand - expanding the tree using the available actions if an unexplored leaf node is selected, (3) rollout - estimating long term reward using a neural network for the action-state pair for the new node added to search tree, and (4) backup - propagating reward value from the leaf nodes upwards to update the evaluations of its ancestors.

Along with this, they propose a new method of adding extra variables in their state representation which helps factor in the parallelization of the scheduled operations, thereby pruning the depth of the output circuit. The way parallel execution works is shown in Figure 3.1, where it is possible to execute multiple gates in one timestep, because the qubits they interact with are not already in a different operation. They referred to their procedure as QRoute, where it performs qubit routing in a hardware agnostic manner, and it outperforms other available qubit routing implementations on various circuit benchmarks.

Yet, they mention that an optimization in runtime can be done through the use of a different kind of neural network approach. QRoute also shares an issue in mitigating by the explicit tree search which takes into account the rewards that will be accrued in the longer-term future. There is scope to further improve this by feeding the entire list of future CNOT-gates directly into their neural network by using transformer encoders to handle the arbitrary length sequence data. On

the whole, the Monte Carlo Tree Search for building up solutions in combinatorial action spaces has exceeded the current state of art methods that perform qubit routing. Despite its success, they note that QRoute is a primitive implementation of their ideas, and there is great scope of improvement in future.

Many different state-of-the-art approaches have been undertaken by researchers. The Monte Carlo Tree Search approach is one of the newer approaches that, in combination with Reinforcement Learning and an initial qubit placement, can lead to remarkable results. This is something that needs to be explored. In the next chapter, the initial qubit placement will be further explained.Many different state-of-the-art approaches have been undertaken by researchers. The Monte Carlo Tree Search approach is one of the newer approaches that, in combination with Reinforcement Learning and an initial qubit placement, can lead to remarkable results. This is something that needs to be explored. The initial qubit placement will be explained in more detail in the next chapter.

Chapter 4

Initial Placement of the Qubits

This section will give answer to the question 'How will an initial qubit placement procedure help with the minimization of SWAP-gates?' by introducing the qubit allocation problem and giving a understanding of some of the possible approaches to solve this problem.

4.1 Qubit Allocation

Given a quantum circuit and an architecture, there must be a way to see if it is possible to map the logical qubits in the former to the physical qubits in the latter. This is called initial qubit placement. The moment that an initial qubit placement ensures that all operations in a circuit are possible, the use of SWAP-gates is no longer necessary. However, it is possible that not all gates can be executed. The moment an allocation is set up in such a way that, for example, 80% of the gates are operable, it may be that fewer SWAP-gates are needed to apply, than when a random allocation is performed. In the paper written by Siraichi et al. [12] a mathematical definition is given for the qubit allocation problem. A simplified explanation would be, suppose there is a topology of the coupling between the physical qubits, of which a list ϕ is made with qubits that interact with each other and a list ψ with qubits involved in a CNOT-gate operations of the circuit. The moment an order of ψ can be arranged in such a way that most or all CNOT-gate operations become possible. The solution given by them has been applied to a Directed Acyclic Graph Topology. Which means the direction of the edges matters. To illustrate this, assume the topology of Figure 4.1.

Figure 4.1: Toy topology for the initial qubit placement

And circuit interacting on this topology illustrated in Figure 4.2.

Figure 4.2: Toy circuit model of CNOT gates for the initial qubit placement

4.2 Weighted Directed Acyclic Graph

To devise a way to arrange the logical qubits in the topology, in such a way that all CNOT-gate operations are possible, it is necessary to look at what kind of topology is expected. This is done by creating a directed acyclic graph based on, based on the out-degree of the qubits. An out-degree is the number of edges going out from a node. So it is counted how many interactions each qubit has to engage in. Using the example of circuit 4.2, the Directed Acyclic Graph 4.3 can be made, illustrating the operations that need to be performed.

Figure 4.3: Coupling Directed Acyclic Graph of the toy circuit model

In the list ψ all CNOT-gate operations are written out in a set with the interacting qubits. What can be inferred from this graph is that when all the out-degrees of every qubit in the graph (the number of arrows) are added up, a logical order of most out-degree would be $Q0 - Q2 - Q1 - Q3$. Another thing that can be made out is the different interactions that each qubit has. For example, it is clear that Q0 undergoes two different interactions, with Q2 and Q3. From this we can already say that Q0 must come to a position where the physical qubit interacts with two other physical qubits. It can also be established that Q3 itself has no out-degrees, so Q3 must be allocated in a position in where there are no out-degrees either.

If the out-degree criterion is used as stimulus to allocate the logical qubits to physical qubits, it becomes possible to make the CNOT-gate operations executable. The table below shows an overview of the number of out-degrees of each node in the topology.

Node	number of out-degrees
A	2
B	1
C	0
D	2
E	1

Qubits	number of out-degrees
Q0	3/2 (3 in total, 2 unique)
Q1	1
Q2	2/1 (2 in total, 1 unique)
Q3	0

When the topology of the physical qubits and the Weighted Directed Acyclic Graph are combined, an allocation like the one shown in Figure 4.4 would yield based on the matching of the number of out-degrees.

Figure 4.4: talk about how we start with the highest out-degree and match them on the highest out-degree of the actual topology, we do this recursively. starting from (a) and going from there.

Because Q0 has the most out-degrees, it is allocated first. Since it only has two unique out-degrees, it is matched with a node that also has two out-degrees. In this case, that would be node A and D. Because the topology maintains a mirror image, a node is randomly chosen to allocate Q0 on. In the case of Figure 4.4(a), it is allocated to node D. The next qubit with the highest out-degree is Q2, which according to the Weighted Directed Acyclic Graph in Figure 4.3 has an edge with Q0. A logical location would be node E (Figure 4.4(b)), because it is connected to Q0 and despite having 2 total out-degrees, it has only 1 unique

out-degree. The qubit allocated next is Q3. The reason that Q1 does not act first is, because Q1 has a directed edge to Q2, but is not reachable. Q3 would be a simple allocation, as it matches the location of node C (Figure 4.4(c)). They both have an out-degree of 0 and do not point anywhere. There are two qubits pointing to Q3, so it makes sense to allocate it in the middle node. Next, the shortest route to Q2 is considered for the allocation of Q1. In this case, the route from node A and B is the same, so the allocation will have to be chosen arbitrarily. In the case of Figure 4.4(d) node A has been chosen. Since Q1 has no direct edge to Q2, this operation is inoperable. In this case, the next step is to use SWAP gates to have the two qubits contiguous, making the circuit executable.

4.3   Weighted Undirected Graph

In the case of an Undirected Graph topology, the direction of each edge no longer matters. In this case, the tie-breaking decision can be made based on the number of edges, independent of the direction. Assume the Undirected Graph topology from Figure 4.5.

Figure 4.5: Undirected Graph Topology

Again, each node has its own number of edges. The table below gives a representation of this.

Node	number of out-degrees
A	2
B	2
C	4
D	2
E	2

As shown in the table, there are many nodes that have the same value of edges. This makes the initial placement procedure difficult. This small example is a simple representation of the real problem. The moment logical qubits have to be allocated in a topology, such as Google's Sycamore device graph (Figure 2.4), things get a lot more difficult. Figure 4.6 shows the Nondeterministic Polynomial time complexity classes. In Computer Science the performance or complexity of an algorithm is described in certain time complexity. The time complexity describes the worst-case scenario, and can be used to describe the execution time required by an algorithm. Easy algorithms run in constant time, but large algorithms for example can quickly be expensive in time. The moment an algorithm

has a polynomial time complexity, this means that the algorithm is heavy in computing power for the computer. Polynomial complexity is a sub category of the Nondeterministic Polynomial (NP) time complexity class, which can be seen in Figure 4.6. The higher up in the intersections, the harder it is for the computer to execute the algorithm. NP-problems can usually still be performed by a computer, but the moment a problem is NP-Complete or NP-Hard, enormously heavy computing power is required from the computer to execute the algorithm.

The problem of initially allocating qubits in a large topology is known as an NP-Complete problem. There appears to be no straight-forward solution for the initial placement procedure problem on large topologies, so that randomness is often still needed to help determine the allocation of the logical qubits.

Figure 4.6: Nondeterministic Polynomial time complexity classes

In Figure 4.7 the same sketch is made of a way in which the qubits can be allocated on the basis of matching the number of edges. In this case it is possible to make the interaction between Q1 and Q2, because the problem of the directions of the edges no longer matters. This ensures that all operations in list ψ (Figure 4.3 can be performed, which means that the circuit in Figure 4.2 is operable without using any SWAP-gates.

Figure 4.7

4.4 Subgraph Isomorphism Partitioning

The process of checking whether two graphs fit together is called Isomorphism in mathematics. Because the topology is often larger than the weighted graph, that arises from a circuit, the problem that is faced is called Subgraph Isomorphism. It is checked whether a part of a large graph matches another smaller graph, see Figure 4.8.

Figure 4.8

Initial qubit placement checks whether the pattern of the subgraph fits into the pattern of the topology graph. The qubit allocation problem can be broken down into two subproblems: (a) Subgraph Isomorphism and (b) Token Swapping problem. Token Swapping is all about finding the shortest route to a target, in this case finding the minimum number of SWAP gates leading to a working circuit.

In a topology where most edges have the same outdegrees, this can cause the solution space to be too large. In most cases a complete weighted graph, like in Figure 4.3, it is unlikely for the graph to be Isomorphic with its topology graph. Sirarchi came up with an update to his approach [13]. A this point an option for this problem can be to look at Operation Partitioning, where CNOT operations in the list are divided into pieces and Subgraphs are made of these operations. By checking these Subgraphs on Isomorphism, an overview is created of the qubits, which are related to each other, that can be mapped directly in the topology. Figure 4.9 shows how this works.
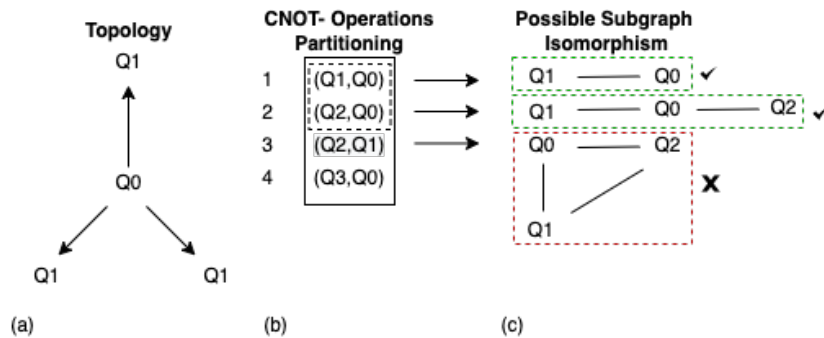
Figure 4.9

Take the topology in Figure 4.9(a) and the list of CNOT operations (b). By making Subgraphs of the operations, the Isomorphism can be looked at. The moment a Subgraph is not Isomorphic with the topology, another mapping has to be looked at and the maximum number of qubits that are Isomorphic with the

topology is remembered. From this point, a mapping can be created from the Subgraph with the most qubits. From this point, it is possible that not all operations satisfied. At this point, the tie-breaker will be the Token Swapping, which will complete the mapping by adding SWAP-gates to make the other CNOT-gates possible. The way the Token Swapping is handled is through the use of the Monte Carlo Tree Search, which is further explained in Section 5.

With these approaches it is possible to ensure that number of SWAP gates are minimized. By devising an allocation that is beneficial for the circuit, it can be ensured that the addition of SWAP-gates does not have to be necessary. The next chapter explains how to add a SWAP gate using the Monte Carlo Tree Search in combination with Reinforcement Learning when it is necessary to add a SWAP-gate.

Chapter 5

Monte Carlo Tree Search with Reinforcement Learning

This chapter will describe how the Monte Carlo Tree Search (MCTS) in combination with Reinforcement Learning (RL) has been applied to optimize quantum circuit routing.

5.1 Implementation of the MCTS

In 2016, the computer program AlphaGo was nominated as the breakthrough of the year. The program consists of a combination of advanced MCTS with deep neural networks and is made to play the Asian game Go [14]. The idea for using the MCTS for the quantum circuit routing is inspired by AlphaGo, where all possible actions are simulated, until a best answer comes out. MCTS is used to determine subsequent actions. To give an illustration of how it works, imagine the game Tic-Tac-Toe, where the best actions for crosses must be be decided 5.1.


Figure 5.1: Simple illustration of simulation of Tic-Tac-Toe. Rewards being add when crosses wins, subtracted when noughts wins and no reward will be given when it is a tie.

To determine the best action, a tree is made consisting of all possible actions, from which all branches arise, each with its own branch. When the game is

over, rewards are given out. When crosses has won, that whole branch gets a
reward of 10. When crosses has lost, the reward is -10 and when no one has won,
nothing will be added or subtracted. MCTS offers a number of advantages over
traditional tree search methods. MCTS does not require any strategic or tactical
knowledge about the given domain to make reasonable decisions. The algorithm
can function effectively with no knowledge of a game apart from its legal moves
and end conditions; this means that a single MCTS implementation can be reused
for a number of games with little modification, and makes MCTS a potential boon
for general game playing. MCTS performs asymmetric tree growth that adapts
to the topology of the search space. The algorithm visits more interesting nodes
more often, and focusses its search time in more relevant parts of the tree. This
makes MCTS suitable for games with large branching factors such as 19x19 Go.
Such large combinatorial spaces typically cause problems for standard depth- or
breadth-based search methods, but the adaptive nature of MCTS means that it
will (eventually) find those moves that appear optimal and focus its search effort
there.

The way in which the MCTS chooses the best action based on the step-by-step
plan shown in figure 5.2.
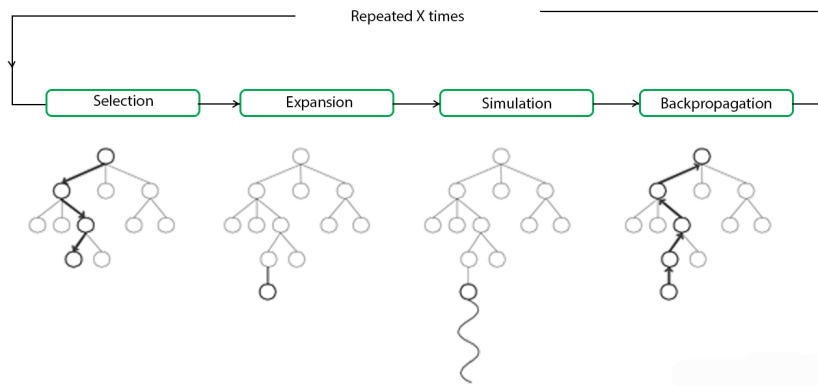
Figure 5.2: Monte Carlo Tree Search roadmap for choosing the best action.

The tree search first starts with a state that by simulating a number of random
nodes. A node represents a possible action. A node can be a parent or a child.
The moment a node is added, the parent and child nodes move down. So the
bottom nodes become the child nodes and the node above it becomes the parent
node. Each node consists of two values, the Upper Confidence Bound applied on
trees (UCT) and a heuristic value that comes from the neural network. UCT is
the factor that decides which node evaluate next in order to maximize probability
of victory from the given state. 5.3.

Figure 5.3: Selection of nodes with Upper Bound Confidence applied on trees (UCT) and a heuristic value H given by the neural network.

5.1.1 Selection

The selection is guided by the UCT, which treats the problem as the multi-armed bandit problem. Imagine you are in a casino facing multiple slot machines and each is configured with an unknown probability of how likely you can get a reward at one play. To figure out the best strategy with highest reward, the exploration vs exploitation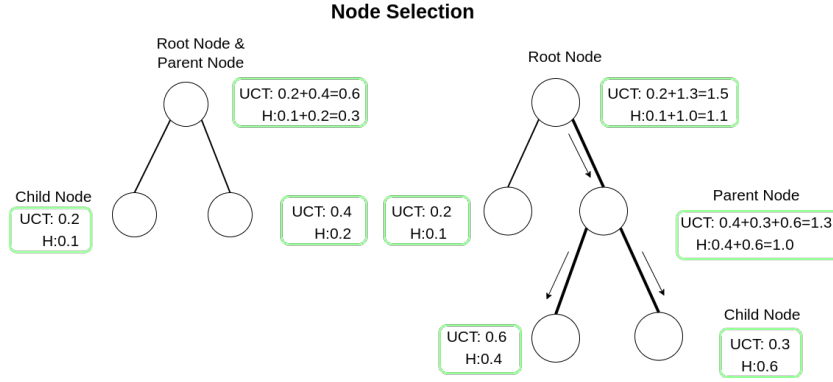 dilemma is demonstrated. Exploitation is the task to select the move that leads to the best results so far. Exploration deals with less promising moves that still have to be examined, due to the uncertainty of the evaluation. Say, one is looking for a new restaurant, it is very likely to eat unpleasant food from time to time at these new restaurants. If all the information about the environment is known, it is possible to find the best restaurant by even just simulating brute-force, so by checking every restaurant. The dilemma comes from the incomplete information, it is necessary to gather enough information to make the best decisions while keeping the risk under control. With exploitation, advantage is taken from the best known option. When exploring, we take some risk to gather information about unknown options. For calculating the UCT, the standard formula is stated as:

$$UCT(node_i) = \bar{x} + c\sqrt{\frac{logN}{n_i}} \tag{5.1}$$

Where $\bar{x}$ denoted the mean node value, N number of times the parent has been visited, n_i number of times the child has been visited and c being a constant.

For the quantum circuit routing optimization, this formula has been built up to the following formula:

$$UCT(node_i) = reward + \sqrt{2}\sqrt{\frac{logN + e + 10^{-6}}{n_i + 10^{-1}}} \tag{5.2}$$

Where $\bar{x}$ is depicted as the reward of that node. The number of times the parent and child have been visited is multiplied by some factors, to keep the answer within a range of 0 and 1. The left term (reward) is the exploitation term. It is simply the average win rate, going larger the better a node has historically performed. The right term $\sqrt{\frac{logN + e + 10^{-6}}{n_i + 10^{-1}}}$ is the exploration term. It goes larger

the less frequently a node is selected for simulation. The exploration parameter c is just a number we can choose that allows us to control how much the equation favors exploration over exploitation; the usual number chosen is $c = \sqrt{2}$ [15].

The UCT is calculated for each node. The tree is simulated up to N number of layers. Suppose N is equal to 3, i.e. as shown in Figure 5.2 the selection part in the tree goes down 3 nodes from the root, with the last nodes representing the leaf nodes. Leaf nodes are the nodes that do not have child nodes themselves. From the leaf node, all UCT are added together to calculate the UCT of the parent nodes. The branch with the highest UCT gets to be picked for the next phase.

To give a feeling of how this phase would work, see the pseudo code below:

Algorithm 1 Selection MCTS

1: **function** $\mathrm{SELECTION}(node)$
2:     **while** node has children **do**
3:         node = child of node with biggest UCT
    return node

5.1.2 Expansion

After the selecting phase, the next phase is the expanding phase. This part is used to increase the options further in the game by expanding the selected node and creating many children nodes. Expansion is the strategic task that, for a given leaf node, decides whether this node will be expanded. The simplest rule is to expand one node per simulated game. If the leaf node is already at the end, the expansion part is not necessary.

Expansion

Figure 5.4: .

Algorithm 2 Expansion MCTS

1: **function** EXPANSION($node$)
2:     **if** node is visited **then**
3:         generate children of node

5.1.3   Simulation

In the simulation phase (also called playout) the best children/ leaf node from the group has to be found. The move which will perform best and lead to the correct answer down the tree is decided here. Through the use of Reinforcement Learning (RL) random decisions can be made in the game further down from every children node. A reward is given to every children node by calculating how close the output of their random decision was from the final output that we need to win the game.


Figure 5.5: .

For example: In the game of Tic-Tac-Toe. If the random decision to make X next to previous X in the game result in three consecutive X-X-X, that means the game has ended. The strategic task that selects moves is through self-play until the end of the game. This task might consist of playing plain random moves or pseudo-random moves chosen according to a simulation strategy. The main idea is to play interesting moves by using patterns, capture considerations, and proximity to the last move. It is the RL model's job to figure out these patterns. When extra rules are added, the condition of limited number of moves is satisfied. When rules as: (1) The game is stopped and scored if it exceeds a given number of moves or (2) the player can only play after the opponent has set a move. Elaborating an efficient simulation strategy is a difficult issue. If the strategy is too stochastic, for instance if it selects moves nearly randomly, then the moves played are often weak. In contrast, if the strategy is too deterministic, for instance if the selected

move for a given position is almost always the same, i.e. too much exploitation takes place, then the exploration of the search space becomes too selective.

The way simulation works is shown in the following pseudo code snippet:

Algorithm 3 Simulation MCTS

1: **function** SIMULATION(*node*)
2:     **while** node state is not final **do**
3:         get next node random
4:         node = next node
    return node state

5.1.4   Backpropagation

Backpropagation is the procedure that propagates the result of a simulated game backwards from the leaf node to the nodes it had to traverse to reach this leaf node. This result is counted positively (reward= +1) if the game is won, and negatively (reward= −1) if the game is lost. Draws lead to a result that has not changed (reward= 0). The value of a node is computed by taking the average of the results of all simulated games made through this node. However, the best results in game play have been obtained by using the plain average of the simulations.

After the simulations, the move finally played by the program in the actual game is the one corresponding to the "best child" of the root. There are different ways to define which child is the best.

- **Max child**: The max child is the child that has the highest value.

- **Robust child**: The robust child is the child with the highest visit count.

- **Robust-max child**: The robust-max child is the child with both the highest visit count and the highest value. If there is no robust-max child at the moment, more simulations are played until a robust-max child occurs.

- **Secure child**: The secure child is the child that maximizes a lower UCT, from all the branches with a winning outcome.

Backpropagation

Figure 5.6: .

To understand the backpropagation, the pseudo code below will summarize how it is programmed.

Algorithm 4 Backpropagation MCTS

1: **function** $\text{BACKPROPAGATION}(node)$
2:      **while** node is not null **do**
3:          increment node score
4:          increment node visits
5:          node = parent node

5.2   Implementation of the Reinforcement Learning Model

The MCTS is aided by the use of Reinforcement Learning (RL). RL is a machine learning training method based on rewarding desired behaviors and/or punishing undesired ones. In general, a reinforcement learning agent is able to perceive and interpret its environment, take actions and learn through trial and error. The RL will be the potential decision-maker in the MCTS. The goal of a RL algorithm is to find a strategy that will generate the optimal outcome. The difference between RL and supervised machine learning, is that the RL model has no labelled data, it has to learn based on experience. The way RL achieves this goal is by allowing a piece of software called an agent to explore, interact with, and learn from the environment. See Figure 5.7.
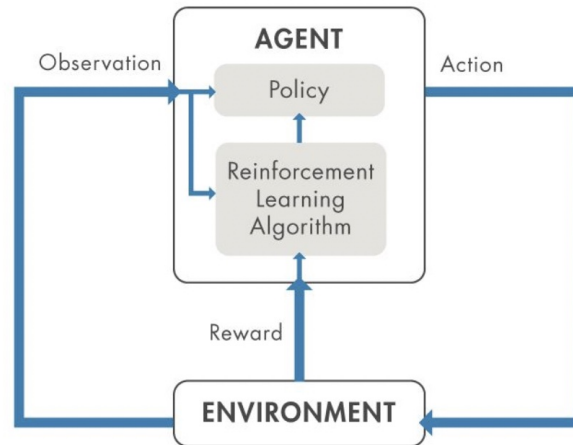
Figure 5.7: Reinforcement learning overview diagram.

Take as an example the maze in Figure 5.8. To learn an RL model to find its way out of a maze, it needs a number of defined attributes: an environment, state (observation), reward and action. In this example, the environment is the maze itself. The state is where the agent is located at the moment. Reward is what the agent receives depending on his action. In the cases of Figure 5.8 a, b and c the rewards will be negative. The moment the agents found its way out of the maze, Figure 5.7 d, the agent will receive a high reward. When the agent still has room left for actions, a positive reward will be given, so that the agent knows that he is still doing well. Based on the rewards, the agent will adjust its policy to come up with a new tactic. The model learns based on a neural network that falls under the block Reinforcement Learning Algorithm in Figure 5.7.


(a) simulation 1   (b) simulation 2   (c) simulation 3   (d) simulation n

Figure 5.8: Example of reinforcement learning solving a maze.

The neural network generates an action based on the state and reward it receives as input. The neural network trains the agent to deliver better results. In the case of the maze, the actions can be to go up, down, left or right. The neural network outputs a probability for each of these actions based on what it things is the best action, see Figure 5.9. Each node, or artificial neuron, connects to another and has an associated weight and threshold. If the output of any individual node is above the specified threshold value, that node is activated, sending data to the next layer of the network. Otherwise, no data is passed along to the next layer of the network.

Figure 5.9: Example of a fully connected neural network.

5.2.1   Neural Network Explained

Think of each individual node as its own linear regression model, composed of input data, weights, a bias (or threshold), and an output. The formula would look something like this:

$$\sum_{i=1}^{m} w_i x_i + bias = w_1 x_1 + w_2 x_2 + w_3 x_3 + bias \tag{5.3}$$

Once an input layer is determined, weights are assigned. These weights help determine the importance of any given variable, with larger ones contributing more significantly to the output compared to other inputs. All inputs are then multiplied by their respective weights and then summed. Afterwards, the output is passed through an activation function, which determines the output. If that output exceeds a given threshold, it "fires" (or activates) the node, passing data to the next layer in the network. This results in the output of one node becoming in the input of the next node. This process of passing data from one layer to the next layer defines this neural network as a feedforward network.

Most deep neural networks are feedforward, meaning they flow in one direction only, from input to output. However, you can also train your model through backpropagation; that is, move in the opposite direction from output to input. Backpropagation allows us to calculate and attribute the error associated with each neuron, allowing us to adjust and fit the parameters of the model(s) appropriately.

Let's break down what one single node might look like using binary values. We can apply this concept to a more tangible example, like whether you should go surfing (Yes: 1, No: 0). The decision to go or not to go is our predicted outcome, or y-hat. Let's assume that there are three factors influencing your decision-making:

- Are the waves good? (Yes: 1, No: 0)

- Is the line-up empty? (Yes: 1, No: 0)

- Has there been a recent shark attack? (Yes: 0, No: 1)

Then, let's assume the following, giving us the following inputs:

- X1 = 1, since the waves are pumping

- X2 = 0, since the crowds are out

- X3 = 1, since there has not been a recent shark attack

Now, we need to assign some weights to determine importance. Larger weights signify that particular variables are of greater importance to the decision or outcome.

- W1 = 5, since large swells do not come around often

- W2 = 2, since one is used to the crowds

- W3 = 4, since one may have a fear of sharks

Lets also assume a threshold value of 3, which would translate to a bias value of –3. With all the various inputs, we can start to plug in values into the formula to get the desired output.

$$Y - hat = (1 * 5) + (0 * 2) + (1 * 4)3 = 6 \tag{5.4}$$

Using the activation function from the beginning of this section, one can determine that the output of this node would be 1, since 6 is greater than 0. In this instance, you would go surfing; but if we adjust the weights or the threshold, we can achieve different outcomes from the model. When we observe one decision, like in the above example, we can see how a neural network could make increasingly complex decisions depending on the output of previous decisions or layers.

In the example above, we used perceptrons to illustrate some of the mathematics at play here, but neural networks leverage sigmoid neurons, which are distinguished by having values between 0 and 1. Since neural networks behave similarly to decision trees, cascading data from one node to another, having x values between 0 and 1 will reduce the impact of any given change of a single variable on the output of any given node, and subsequently, the output of the neural network.
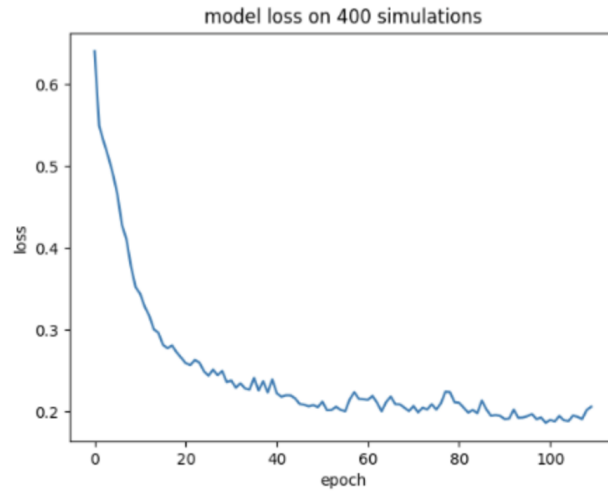
Chapter 6

Results

This chapter discusses the results of the MCTS..

6.1   Results of the MCTS

Model was trained based on random qubit allocations and random circuits.

Figure 6.1: Model loss after training on 400 MCTS simulations

Figure 6.2: Random Qubit Allocation in Linear Topology

During the MCTS process, the circuit in Figure 6.3 has been used as an example with Figure 6.5 as the topology. Circuit has been encoded as:

$$[[0, 1, 0], [3, 2, 0], [3, 0, 0], [0, 2, 0], [1, 2, 0], [1, 0, 0], [2, 3, 0], [3, 0, 0], [0, 2, 0]] \tag{6.1}$$

Various tests have been done with this circuit, in which a random allocation has been done. The results of this can be seen in Figures 6.4.

Figure 6.3: Inoperable Test Circuit

Figure 6.4: Circuit optimized with random qubit allocation

These figures show that the SWAP-gates ensure that the circuits are compatible with respect to the topology. Looking at the results using circuits with different lengths, Figure 6.7 shows the results based on the difference in depth of those circuits.

Figure 6.5: Circuit Depth Results (still need to write a caption)

6.2 Benchmarking of the Circuit Depth

figures/qiskit_depth.png

Figure 6.6: Mine, Qiskit, TKet, Pozzi circuit depth

Figure 6.7: processor time with increase in number of qubits

Chapter 7

Discussion

In this thesis research is done on quantum circuit routing optimization by asking the main question 'How can quantum circuit routing be optimized using Reinforcement Learning applied on Monte Carlo Tree Search?'. In order to answer this question, research has been done into why there are issues with quantum circuit routing. Another thing that has been looked at is what state-of-the-art approaches already exist to solve this problem. Research has been done in how initial qubit placement can provide a minimized number of SWAP gates and finally, two state-of-the-art approaches were combined to see how they can help optimize quantum circuit routing.

7.1 Key Findings of the Results

From this research, the results suggest that ... This supports the theory... The analysis identifies...

7.1.1 Interpretation of the Results

The results relate to the research that has been done, based on.. The results differ from the previous studies, because.. The way these results contribute to the field of Quantum Computing is ...

7.2 Limitation of the Study

During the research, some limitations were detected. ...

7.3 Recommendation and Future Work

Chapter 8

Conclusion

Bibliography

[1] B. Tan and J. Cong, "Optimality study of existing quantum computing layout synthesis tools," *IEEE Transactions on Computers*, vol. 70, pp. 1363–1373, 2021. (document), 2.4

[2] "Quantum computing transpiler documentation," May 2021. [Online]. Available: https://qiskit.org/documentation/apidoc/transpiler.html (document), 3.2

[3] G. Li, Y. Ding, and Y. Xie, "Tackling the qubit mapping problem for nisq ... - arxiv.org," May 2019. [Online]. Available: https://arxiv.org/pdf/1809.02573.pdf (document), 3.2, 3.3

[4] F. Bova, A. Goldfarb, and R. Melko, "Quantum computing is coming. what can it do?" Jul 2021. [Online]. Available: https://hbr.org/2021/07/quantum-computing-is-coming-what-can-it-do 1.1

[5] M. G. Pozzi, S. J. Herbert, A. Sengupta, and R. D. Mullins, "Using reinforcement learning to perform qubit routing in quantum compilers," Jul 2020. [Online]. Available: https://arxiv.org/abs/2007.15957 1.4, 3.3

[6] A. Sinha, U. Azad, and H. Singh, "Qubit routing using graph neural network aided monte carlo tree search," Mar 2022. [Online]. Available: https://arxiv.org/abs/2104.01992 1.4, 3.4

[7] C. Gronlund, M. Li, and S. Lopez, "Dirac-notatie - azure quantum," Apr 2022. [Online]. Available: https://docs.microsoft.com/nl-nl/azure/quantum/concepts-dirac-notation 2.2

[8] D. Shaw, "Quantum hardware - into the quantum jungle," Jul 2020. [Online]. Available: https://www.factbasedinsight.com/quantum-hardware-into-the-quantum-jungle/ 2.4

[9] "Cambridge quantum releases tket version 0.7 with open access to all python users," Nov 2021. [Online]. Available: https://cambridgequantum.com/cambridge-quantum-releases-tket-v0-7/ 3.1

[10] A. Cowtan, S. Dilkes, R. Duncan, A. Krajenbrink, W. Simmons, and S. Sivarajah, "On the qubit routing problem - dagstuhl," Feb 2019. [Online]. Available: https://drops.dagstuhl.de/opus/volltexte/2019/10397/pdf/LIPIcs-TQC-2019-5.pdf 3.1

[11] S. Herbert and A. Sengupta, "Using reinforcement learning to find efficient qubit routing policies for deployment in near-term quantum computers," 12 2018. 3.3

[12] M. Siraichi, V. F. Dos Santos, C. Collange, and F. M. Quintão Pereira, "Qubit allocation," Feb 2018. [Online]. Available: https://hal.archives-ouvertes.fr/hal-01655951/file/Siraichi_QubitAllocation_CGO18.pdf 4.1

[13] M. Y. Siraichi, V. Fernandes Dos Santos, C. Collange, and F. Magno Quintão Pereira, "Qubit allocation as a combination of subgraph isomorphism and token swapping," Oct 2019. [Online]. Available: https://homepages.dcc.ufmg.br/~fernando/publications/papers/OOPSLA19.pdf 4.4

[14] "Alphago." [Online]. Available: https://www.deepmind.com/research/highlighted-research/alphago 5.1

[15] M. Lu, "General game-playing with monte carlo tree search," Oct 2017. [Online]. Available: https://medium.com/@quasimik/monte-carlo-tree-search-applied-to-letterpress-34f41c86e238 5.1.1

[16] J. Preskill, "Quantum Computing in the NISQ era and beyond," *Quantum*, vol. 2, p. 79, Aug. 2018. [Online]. Available: https://doi.org/10.22331/q-2018-08-06-79

[17] F. Admin, "Qubit dashboard," Jul 2020. [Online]. Available: https://www.factbasedinsight.com/qubit-dashboard/

[18] J. Hui, "Monte carlo tree search (mcts) in alphago zero," May 2018. [Online]. Available: https://jonathan-hui.medium.com/monte-carlo-tree-search-mcts-in-alphago-zero-8a403588276a

[19] D. Silver, "Alphago zero: Starting from scratch," Oct 2017. [Online]. Available: https://www.deepmind.com/blog/alphago-zero-starting-from-scratch

[20] G. Nannicini, L. S. Bishop, O. Gunluk, and P. Jurcevic, "Optimal qubit assignment and routing via integer programming," Jul 2021. [Online]. Available: https://arxiv.org/abs/2106.06446

[21] "Cambridge quantum computing announces update to t: Ket quantum software development kit," Oct 2020. [Online]. Available: https://www.hpcwire.com/off-the-wire/cambridge-quantum-computing-announces-update-to-tket%E2%9F%A9-quantum-software-development-kit/

[22] A. Cotarelo, V. García-Díaz, E. R. Núñez-Valdez, C. González García, A. Gómez, and J. Chun-Wei Lin, "Improving monte carlo tree search with artificial neural networks without heuristics," Feb 2021. [Online]. Available: https://mdpi-res.com/d_attachment/applsci/applsci-11-02056/article_deploy/applsci-11-02056-v2.pdf?version=1614749006

[23] I. C. Education, "What are neural networks?" [Online]. Available: https://www.ibm.com/cloud/learn/neural-networks